

CS425 - Project 3 Report

Nishant Gupta, nishgu@iitk.ac.in, 13447

September 30, 2016

STCP Protocol

1 Implementation Specs

- Sliding window is implemented
- Proper endianness handling is insured
- Abrupt disconnects are handled
- Proper errno is set

2 Test results

- All Combinations of two clients(reference client and client.c in skeleton) and two server(reference server and server.c in skeleton) are working properly
- The client exits cleanly with **Ctrl+D** and **Ctrl+C** interrupts and the server with **Ctrl+C** interrupt.
- Single file transfer with **-f** flag is working properly
- Multiple client connections are working fine with one client queued after other

3 Appendix

3.1 Source code :

```
/*
 * transport.c
 *
 * Project 3
 *
 * This file implements the STCP layer that sits between the
 * mysocket and network layers. You are required to fill in the STCP
 * functionality in this file.
 */

#include "transport.h"
#include "mysock.h"
#include "stcp_api.h"
#include <arpa/inet.h>
#include <assert.h>
#include <stdarg.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <errno.h>

enum { CSTATE_ESTABLISHED }; /* you should have more states */

/* this structure is global to a mysocket descriptor */
typedef struct {
    bool_t done; /* TRUE once connection is closed */

    int connection_state; /* state of the connection (established, etc.) */
    tcp_seq initial_sequence_num;
    tcp_seq initial_sequence_num_peer;
    tcp_seq current_sequence_num_peer;
    tcp_seq current_sequence_num;
```

```

    tcp_seq sender_window, unacked, congestion_window,
        last_acked;
    /* any other connection-wide global variables go here */
} context_t;

static void generate_initial_seq_num(context_t *ctx);
static void control_loop(mysocket_t sd, context_t *ctx);

inline void hton(STCPHeader* header){
    header->th_seq = htonl(header->th_seq);
    header->th_ack = htonl(header->th_ack);
    header->th_win = htons(header->th_win);
}

inline void ntoh(STCPHeader* header){
    header->th_seq = ntohl(header->th_seq);
    header->th_ack = ntohl(header->th_ack);
    header->th_win = ntohs(header->th_win);
}

/* initialise the transport layer, and start the main loop, handling
 * any data from the peer or the application.  this function should not
 * return until the connection is closed.
 */
void transport_init(mysocket_t sd, bool_t is_active) {
    context_t *ctx;
    int bytes_recieved, send_status;
    ctx = (context_t *)calloc(1, sizeof(context_t));
    assert(ctx);

    generate_initial_seq_num(ctx);
    STCPHeader *hdr = (STCPHeader *)calloc(1, sizeof(STCPHeader));
    /* XXX: you should send a SYN packet here if is_active, or wait for one
     * to arrive if !is_active.  after the handshake completes, unblock the
     * application with stcp_unblock_application(sd).  you may also use
     * this to communicate an error condition back to the application, e.g.
     * if connection fails; to do so, just set errno appropriately (e.g. to
     * ECONNREFUSED, etc.) before calling the function.
     */
    if (is_active) {
        /**
        printf("Active Connection initiation request\n");
        */

        // send SYN packet
        hdr->th_flags = TH_SYN;
        hdr->th_seq = ctx->current_sequence_num++;
        hdr->th_off = sizeof(STCPHeader)/4;

        hton(hdr);
        send_status = stcp_network_send(sd, hdr, sizeof(STCPHeader), NULL);
        if(send_status == -1) {
            errno = ECONNREFUSED;
            return;
        }
    }
}

```

```

// recv and process SYN+ACK packet
bytes_recieved = stcp_network_recv(sd, hdr, sizeof(STCPHeader));
if(bytes_recieved == 0) {
    errno = ECONNREFUSED;
    return;
}
ntoh(hdr);
assert(hdr->th_flags == (TH_SYN | TH_ACK));
assert(ctx->current_sequence_num == hdr->th_ack);
ctx->initial_sequence_num_peer = hdr->th_seq;
ctx->current_sequence_num_peer = ctx->initial_sequence_num_peer + 1;

// send ACK packet
memset(hdr, 0, sizeof(STCPHeader));
hdr->th_flags = TH_ACK;
hdr->th_seq = ctx->current_sequence_num;
hdr->th_off = sizeof(STCPHeader)/4;
hdr->th_ack = ctx->current_sequence_num_peer;
hton(hdr);
send_status = stcp_network_send(sd, hdr, sizeof(STCPHeader), NULL);
if(send_status == -1) {
    errno = ECONNABORTED;
    return;
}

// Handshake complete
} else {
    /**
    printf("Passive Connection initiation request\n");
    */
    // recieve SYN packet
    bytes_recieved = stcp_network_recv(sd, hdr, sizeof(STCPHeader));
    if(bytes_recieved == 0) {
        errno = ECONNABORTED;
        return;
    }
    ntohs(hdr);
    ctx->initial_sequence_num_peer = hdr->th_seq;
    ctx->current_sequence_num_peer = ctx->initial_sequence_num_peer + 1;

    // send SYN+ACK packet
    memset(hdr, 0, sizeof(STCPHeader));
    hdr->th_flags = TH_ACK | TH_SYN;
    hdr->th_seq = ctx->current_sequence_num++;
    hdr->th_off = sizeof(STCPHeader)/4;
    hdr->th_ack = ctx->current_sequence_num_peer;
    hton(hdr);
    send_status = stcp_network_send(sd, hdr, sizeof(STCPHeader), NULL);
    if(send_status == -1) {
        errno = ECONNABORTED;
        return;
    }
}

```

```

    // recieve ACK
    stcp_network_recv(sd, hdr, sizeof(STCPHeader));
    if(bytes_recieved == 0) {
        errno = ECONNABORTED;
        return;
    }
    ntohs(hdr);

    assert(hdr->th_flags & TH_ACK);
    assert(ctx->current_sequence_num == hdr->th_ack);
}
ctx->connection_state = CSTATE_ESTABLISHED;
/**/
//printf("Handshake Completed\n"); /**/
stcp_unblock_application(sd);
control_loop(sd, ctx);

/* do any cleanup here */
free(ctx);
}

/* generate random initial sequence number for an STCP connection */
static void generate_initial_seq_num(context_t *ctx) {
    assert(ctx);

#ifdef FIXED_INITNUM
    /* please don't change this! */
    ctx->initial_sequence_num = 1;
#else
    /* you have to fill this up */
    // ctx->initial_sequence_num = 1; // testing

    ctx->initial_sequence_num = rand() % 256;
    ctx->current_sequence_num = ctx->initial_sequence_num;
#endif
}

/* control_loop() is the main STCP loop; it repeatedly waits for one of the
 * following to happen:
 * - incoming data from the peer
 * - new data from the application (via mywrite())
 * - the socket to be closed (via myclose())
 * - a timeout
 */
static void control_loop(mysocket_t sd, context_t *ctx) {
    assert(ctx);
    assert(!ctx->done);
    unsigned int event;
    bool_t this_end_closed = FALSE, other_end_closed = FALSE;
    STCPHeader *hdr = (STCPHeader *)calloc(1, sizeof(STCPHeader));
    int off = sizeof(STCPHeader) / 4 + ((sizeof(STCPHeader) % 4) ? 1 : 0);
    ctx->unacked = 0;
    ctx->sender_window = 3072;
    ctx->congestion_window = 3072;

```

```

int send_status;
void *buffer = (void *)malloc(STCP_MSS + 4 * off);
while (!ctx->done) {

    event = stcp_wait_for_event(sd, ANY_EVENT, NULL);

    if ((event & APP_DATA) && !this_end_closed) {

        int data_limit = MIN(STCP_MSS, ctx->sender_window - ctx->unacked);

        hdr->th_seq = ctx->current_sequence_num;
        hdr->th_off = off;
        hdr->th_flags = 0;
        hdr->th_win = ctx->congestion_window;

        int bytes_to_send = stcp_app_rcv(sd, buffer, data_limit);
        hton(hdr);
        send_status = stcp_network_send(sd, hdr, sizeof(STCPHeader), buffer, bytes_to_send,
                                         NULL);
        if(send_status == -1) {
            errno = ECONNABORTED;
            return;
        }
        ctx->unacked += bytes_to_send;
        ctx->current_sequence_num += bytes_to_send;
    }

    if (event & NETWORK_DATA) {
        int bytes_received = stcp_network_rcv(sd, buffer, STCP_MSS + off * 4);
        if(bytes_received == 0) {
            errno = ECONNABORTED;
            ctx->done = TRUE;
            continue;
        }

        memcpy((void *)hdr, buffer, sizeof(STCPHeader));
        ntoh(hdr);
        ctx->sender_window = MIN(hdr->th_win, ctx->congestion_window);

        if ((hdr->th_flags & TH_ACK) && ctx->last_acked < hdr->th_ack - 1) {
            ctx->last_acked = hdr->th_ack - 1;
            ctx->unacked = ctx->current_sequence_num - hdr->th_ack;
            ctx->sender_window = MIN(hdr->th_win, ctx->congestion_window);
            if(this_end_closed) { // we won't send any data after this
                                // so this should hold
                assert(hdr->th_ack <= ctx->current_sequence_num + 1);
            }
        }

        if (bytes_received - off * 4 > 0) {
            if(ctx->current_sequence_num_peer > hdr->th_seq + bytes_received - off * 4){
                // data is duplicate but send new ack
                hdr->th_ack = hdr->th_seq + bytes_received - off * 4;
                hdr->th_flags = TH_ACK;
                hdr->th_win = ctx->congestion_window;
            }
        }
    }
}

```

```

    hdr->th_off = off;
    hdr->th_seq = ctx->current_sequence_num;

    hton(hdr);
    send_status = stcp_network_send(sd, hdr, off * 4, NULL);
} else {
    int new_data_offset = 0;
    new_data_offset = ctx->current_sequence_num_peer - hdr->th_seq;
    ctx->current_sequence_num_peer = hdr->th_seq + bytes_received - off * 4;
    hdr->th_ack = ctx->current_sequence_num_peer;
    hdr->th_flags = TH_ACK;
    hdr->th_win = ctx->congestion_window;
    hdr->th_off = off;
    hdr->th_seq = ctx->current_sequence_num;

    hton(hdr);
    send_status = stcp_network_send(sd, hdr, off * 4, NULL);
    stcp_app_send(sd, (char *)buffer + off * 4 + new_data_offset, bytes_received - off * 4 - new_
}
if(send_status == -1) {
    errno = ECONNABORTED;
    return;
}
}

if (hdr->th_flags & TH_FIN) {
    ctx->current_sequence_num_peer++; // FIN takes one Segment space
    //ACK for FIN
    memset(hdr, 0, sizeof(STCPHeader));
    hdr->th_ack = ctx->current_sequence_num_peer;
    hdr->th_seq = ctx->current_sequence_num;
    hdr->th_flags = TH_ACK;
    hdr->th_win = ctx->congestion_window;
    hdr->th_off = off;

    hton(hdr);
    send_status = stcp_network_send(sd, hdr, off * 4, NULL);
    if(send_status == -1) {
        errno = ECONNABORTED;
        return;
    }
    other_end_closed = TRUE;
    stcp_fin_received(sd);
}

}

if(event & APP_CLOSE_REQUESTED) {
    //printf("App close requested\n");
    // send FIN
    memset(hdr, 0, sizeof(STCPHeader));
    hdr->th_seq = ctx->current_sequence_num++;
    hdr->th_flags = TH_FIN;

```

```

    hdr->th_off = sizeof(STCPHeader)/4;
    hdr->th_win = ctx->congestion_window;

    hton(hdr);
    send_status = stcp_network_send(sd, hdr, off * 4, NULL);
    if(send_status == -1) {
        errno = ECONNABORTED;
        return;
    }

    this_end_closed = TRUE;
}
if (event & TIMEOUT) {
    ctx->done = TRUE;
}

if(this_end_closed && other_end_closed) ctx->done = TRUE;
/* etc. */
}
free(hdr);
free(buffer);
}

/*****
/* our_dprintf
*
* Send a formatted message to stdout.
*
* format          A printf-style format string.
*
* This function is equivalent to a printf, but may be
* changed to log errors to a file if desired.
*
* Calls to this function are generated by the dprintf amd
* perror macros in transport.h
*/
void our_dprintf(const char *format, ...) {
    va_list argptr;
    char buffer[1024];

    assert(format);
    va_start(argptr, format);
    vsnprintf(buffer, sizeof(buffer), format, argptr);
    va_end(argptr);
    fputs(buffer, stdout);
    fflush(stdout);
}

```